

RAG Pipeline

Model Comparison Report

Generated: 24 May 2026, 03:48 PM · Models evaluated: 1

Pipeline Configuration

Configuration	
Retrieval mode	semantic
Reranker	OFF
Adaptive weights	OFF
Fusion strategy	N/A
BM25 / Sem weights	N/A
Filtering	OFF
Cross-modal boost	ON
text_k / rerank_k / image_k	10 / 5 / 5
Embed model	BAAI/bge-large-en-v1.5
Image embed model	ViT-L-14 (laion2b_s32b_b82k)
LLM model(s)	gemma4:e2b
Chunk tokens / overlap	384 / 128

Metric Comparison — All Models

■■ RETRIEVAL (shared across all models) ■■	
M1 Embedding time ↓	12.6674 s
M2 Index size	176 vectors
M3 Retrieval latency ↓	0.0632 s
M4 Cosine Sim (Text)	0.7503
M4 Cosine Sim (Image)	0.4110
M5 Page Coverage@k ↑	90.00 %
M9 Context length	3308 ch
■■ GENERATION QUALITY (per model, vs ground truth) ■■	
M6 ROUGE-1 ↑	0.5153
M7 ROUGE-2 ↑	0.1976
M8 ROUGE-L ↑	0.2878
M10 BLEU ↑	0.1307
M11 METEOR ↑	0.3353
M12 BERTScore ↑	0.8763
M13 FCD ↓	17.87
M14 Faithfulness ↑	87.86 %
M15 GT Coverage ↑	89.77 %
■■ PERFORMANCE (per model) ■■	
M16 E2E Latency ↓	16.4334 s
M17 Throughput ↑	0.061 q/s
M18 CPU Usage ↓	0.29 %
M19 RAM Usage ↓	0.125 GB
GPU avg	27.48 %
GPU peak	100.00 %

* Retrieval metrics are computed once and are identical across model columns because retrieval runs before any LLM is invoked.