

RAG Pipeline

Model Comparison Report

Generated: 24 May 2026, 04:08 PM · Models evaluated: 1

Pipeline Configuration

Configuration	
Retrieval mode	hybrid
Reranker	OFF
Adaptive weights	OFF
Fusion strategy	RRF
BM25 / Sem weights	0.5 / 0.5
Filtering	OFF
Cross-modal boost	ON
text_k / rerank_k / image_k	10 / 5 / 5
Embed model	BAAI/bge-large-en-v1.5
Image embed model	ViT-L-14 (laion2b_s32b_b82k)
LLM model(s)	gemma4:e2b
Chunk tokens / overlap	384 / 128

Metric Comparison — All Models

■■ RETRIEVAL (shared across all models) ■■	
M1 Embedding time ↓	11.6289 s
M2 Index size	176 vectors
M3 Retrieval latency ↓	0.0484 s
M4 Fused Relevance (Text)	0.7471
M4 Fused Relevance (Image)	0.4108
M5 Page Coverage@k ↑	92.50 %
M9 Context length	3382 ch
■■ GENERATION QUALITY (per model, vs ground truth) ■■	
M6 ROUGE-1 ↑	0.5120
M7 ROUGE-2 ↑	0.1998
M8 ROUGE-L ↑	0.2938
M10 BLEU ↑	0.1274
M11 METEOR ↑	0.3304
M12 BERTScore ↑	0.8777
M13 FCD ↓	18.29
M14 Faithfulness ↑	95.83 %
M15 GT Coverage ↑	91.90 %
■■ PERFORMANCE (per model) ■■	
M16 E2E Latency ↓	16.0490 s
M17 Throughput ↑	0.062 q/s
M18 CPU Usage ↓	0.31 %
M19 RAM Usage ↓	0.125 GB
GPU avg	27.09 %
GPU peak	100.00 %

* Retrieval metrics are computed once and are identical across model columns because retrieval runs before any LLM is invoked.